

What Deception Probes Read and How Evaluations of Them Fail

Andres Garcia
Independent Researcher

Abstract

Linear probes on the residual activations of large language models are the leading proposal for monitoring deceptive generation. We ask what such probes read, where that signal is computed, and whether the standard experiments used to answer these questions are themselves sound. Under a fully pre-specified pipeline: statement-grouped splits, nested layer selection, threshold-free metrics, a single device and precision, on Qwen2.5-0.5B/1.5B/3B and Gemma-2-2b, per-layer ablation curves with a three-label specificity decomposition show that deception-specific damage has no canonical shape: it is band-localized at 1.5B (L15-L17, paired asymmetry to -0.26), multi-site at Gemma-2b (four layers, max -0.31), undetectable at the deployed readout at 3B ($\text{AUC} \geq 0.93$ under every single-MLP ablation), and absent at 0.5B, where the probe is too weak to localize anything. Answer polarity survives every one of 43 ablation conditions (AUC consistently at 1.00), and clean patching restores 95-98% of what ablation removes. The critical layer is re-selected identically under all three prompt formats in both models with localizable structure, while steering resistance is format-dependent (collapse- α grows from 4 to 24). Out of format, instructed-transfer evaluation collapses: naive transfer AUCs span 0.030-1.000 across nine models, in one case exceeding the positive-control ceiling, a naive de-confound statistic spans 0.002-0.994, denial bias leaves seven of nine models unidentifiable (all five Qwen models lie only by denying truths), and in the two identifiable models the instruction coefficient dominates a small, sign-inconsistent deception coefficient. We distill six controls into an evaluation protocol; code, data, splits, and a regenerable results file are released.

1. Introduction

Linear probes on residual activations are the leading proposal for monitoring deception in language models: they are cheap, they report near-ceiling accuracy on their training distributions, and they are already being evaluated for deployment. Three questions underneath that proposal remain open. What do such probes actually read? Where is that signal computed? And are the standard experiments used to answer the first two questions, point interventions at selected components and instructed-lying transfer tests, themselves sound? We answer all three on deliberately small models (0.5B–3B parameters, two families), trading construct realism for full causal access, under a pipeline in which every design choice was pre-specified and every number regenerates from a released results file. To summarize our results, the probes read truth-consistency information; where that information is computed has no canonical shape across models; and the standard evaluation practices returned a wrong answer at nearly every step until controlled.

A single measurement illustrates the problem. Ablating Qwen-3B’s critical MLP drops a probe’s accuracy from 0.950 to 0.719, a fall of 0.231 that reads unambiguously as a necessary component. However, that same ablation only moves the AUC from 0.992 to 0.979, a fall of 0.013.

Table 1: De-calibration exhibit for Qwen-3B highlighting differences in selected measurements

Metric (Qwen-3B)	Clean	Ablated	Difference
Accuracy	0.950	0.719	-0.231
AUC	0.992	0.979	-0.013

Since AUC depends only on the ranking of scores, the representation the probe reads is intact, but what moved was the decision threshold. If we look at accuracy, the field's common choice, this layer is necessary. If we look at a threshold-free metric, this layer is no longer necessary. We found three such reversals in our own results prior to understanding the pattern, and §6 is the set of controls that identified it.

Contributions:

1. **Causal topology, curve-adjudicated (§4):** Band / multi-site / dissolved / absent across four models; polarity intact in every causal run; conduit sufficiency 95-98%; maximum deployed damage from any single MLP ablation 0.32 / 0.15 / 0.06 (across 1.5B, Gemma-2b and 3B, with 0.5B's 0.24 arising solely from readout-self-ablation), a per-model contrast, not a law (pre-registered falsification, App. G).
2. **A methods result (§4, §6):** Magnitude-based point interventions misidentify the mechanism in every model; 1.5B's strongest deployed disruptor has zero deception-specificity; 0.5B's selected "critical" layer is the readout itself. Only per-layer curves with label decomposition adjudicate.
3. **Transfer evaluation fails measurably (§6):** Behavioral degeneracy (denial bias), naive statistics spanning the full [0,1] interval, and instruction-dominated regressions with sign-inconsistent deception loading in the only identifiable models, each caught by a specific control, distilled into a six-step protocol.
4. **A named positive (§3):** Within-format source-generality, probes trained on AI-generated statements transfer to a secondary source matched to the same format (0.75 / 0.92 / 0.96 / 0.955 at the nested layer).

Across four models we ablated every MLP in the recovery band and measured, threshold-free, damage to deception, truth, and polarity readouts. Three regularities survive. Polarity is untouched (AUC maintains 1.000 score) in all 43 ablation conditions (39 per-layer cells plus four paraphrase replications) so no ablated component is a generic bottleneck. Deception-specific damage (paired truth, deception asymmetry, grouped bootstrap) has no clearly defined shape: a contiguous band at Qwen-1.5B (L15-L17: -0.261, -0.194, -0.240), four distinct sites at Gemma-2b (L9/L11/L13/L14, maximal -0.314 at L14, a layer no magnitude criterion selects), weak-probe mid-band sites only at Qwen-3B (whose deployed readout survives every ablation at $AUC \geq 0.934$), and nothing localizable at Qwen-0.5B, where the selection rule degenerates to the readout layer itself. Where identifiable, the components are conduits: clean-patching restores 97% / 98% / 95% of what ablation removes (1.5B / Gemma-2b / 3B; 0.5B has a wide-CI of 0.67). We claim the probes read truth-consistency information whose computation varies in localizability across models, and that every point-intervention summary we computed before the curves, under our own pre-registered rules, misidentified this topology.

Each of the findings above exists in its present form because an earlier version of it was an artifact. Reporting causal effects as detection rate on deceptive items, recall on positives (the field's current default), read decision-boundary shifts as signal in either direction: our strongest necessity effect appeared, under recall, as an ablation that *improved* detection. Balanced accuracy fixed the class asymmetry but still conflated boundary shift with signal destruction: at 3B it manufactured a "critical

MLP" whose ablation leaves the representation fully intact, the accuracy decreased from 0.950 to 0.719 while the AUC stayed roughly the same from 0.992 to 0.979. A restoration formula missing one experimental condition reversed our sufficiency verdict, the corrected conduit restorations are 0.970 / 0.980 / 0.945 across Qwen-1.5B / Gemma-2b / Qwen-3B. A control layer selected under a degenerate rule failed in a way that looked like a discovery, and a clean-room rerun on the same hardware flipped one regression coefficient's sign and dissolved one of our catalogued pitfalls, the preprocessing flip existed only in interaction with the readout-position error (App. E). None of these are exotic mistakes; they are default measurement choices. The protocol of §6 is the set of controls that caught them.

2. Related Work

Deception and lie-detection probes: Linear probes on residual activations are the leading proposal for monitoring deceptive generation: Azaria and Mitchell (2023) train truthfulness probes on internal states; Goldowsky-Dill et al. (2025) report near-ceiling detection on contrasting honesty instructions. Our results are compatible with theirs and locate the difficulty elsewhere: we show that the same instruction contrast that makes such probes trainable is also what they read out of format, and that the evaluations used to certify transfer are unidentifiable in seven of nine models. Pacchiardi et al. (2023) detect lying from follow-up behavior; representation-engineering approaches (Zou et al., 2023) generalize the recipe. Levinstein and Herrmann (2023) are the closest precedent for our transfer results, arguing from generalization failures that such probes track truth-adjacent features rather than lying; we give that claim a causal grounding and show the evaluations themselves are the harder problem.

Geometry of truth: Marks and Tegmark (2023) characterize low-dimensional truth subspaces; Kumar (2026) pressure-tests deception-probe geometry observationally on Gemma-3, attributing fragility to training-distribution narrowness. We are complementary: where that line infers structure from probe behavior, we intervene and find the computation's topology is model-dependent in ways observational geometry cannot distinguish.

Probing methodology: Our protocol extends a tradition of probing skepticism: control tasks (Hewitt & Liang, 2019), amnesic probing (Elazar et al., 2021), the interpretability-illusion line (Bolukbasi et al., 2021), and Belinkov's survey (2021). We contribute the transfer-evaluation analogues of behavioral-degeneracy gates, instruction-leak de-confounds, and position/metric controls, each of which reversed a conclusion in our own study.

Causal interventions: We use standard tools such as causal tracing (Meng et al., 2022) and path-level analysis (Wang et al., 2022), and inherit known cautions on patching metrics (Heimersheim & Nanda, 2024; Zhang & Nanda, 2024), which our recall-inversion results sharpen into a worked failure case. The deployed-readout compensation we observe connects directly to self-repair phenomena, the Hydra effect (McGrath et al., 2023; Rushing & Nanda, 2024): our contribution is showing repair magnitude is model-dependent and can fully absorb single-module ablation.

Evasion: Learned attacks defeat latent-space monitors (Bailey et al., 2024; Gupta & Jenner, 2025; McGuinness et al., 2025). We show measurement breaks before any adversary arrives: white-box steering is our upper bound, and even un-attacked evaluations mislead without the §6 controls.

3. Setup

Pipeline: We started with a set of 400 true or false statements, initially generated by Claude (Opus 4.8 High, May 2026 version), and a secondary set of 100 statements, balanced 50 true and 50 false, written by a non-author with assistance from another AI model (ChatGPT, default model, June 2026). A replication

on an external corpus (Azaria and Mitchell’s true and false statement set) is in progress at probe level only, covering within-corpus accuracy, cross-corpus transfer in both directions, and topic-held-out splits. It is not included in this version, and no claim here depends on it. These statements were separated into statement-grouped splits (both members of a statement pair on one side of every split). We selected the readout layer by nested selection: layer accuracy was averaged over 10 validation splits, and the earliest layer within 0.02 of the peak was chosen, with the held-out test fold never used for the selection; Critical MLPs were pre-registered as the layer whose ablation most reduces the probe’s continuous score on training-fold items. All causal effects reported threshold-free (AUC) with balanced accuracy in the appendix; grouped bootstrap over statements for every interval. One device and precision (A100, fp16); every number regenerates from the released MEASURED.md; per-run manifests in the repository.

Dataset used: The primary dataset is 400 short statements, balanced 200 true and 200 false, each at most twelve words. It was assembled in two batches: 200 statements covering geography, science, history, and elementary school mathematics, and 200 constructed as minimal true and false pairs over the same domains. Statements were drafted with LLM assistance and reviewed by the author for factual accuracy and length. Splits are grouped by statement, so the honest and deceptive presentations of a given statement fall on the same side. Eight statements appear in both batches in near-identical form and are therefore assigned to different groups; two of those pairs fall across a train and test boundary, bounding any inflation of test accuracy at 1.25 percentage points (App. C). An additional set of 100 short statements, balanced 50 true and 50 false, was written by a non-author with assistance from another AI model (ChatGPT, default model, June 2026), delivered as a Word document and converted to JSON by the author.

Label definitions and incoherence handling: Our experiments use two pipelines with different label provenance. In the *causal* pipeline (all ablation, patching, and specificity results), the answer token is teacher-forced: each statement is presented with an appended Yes/No answer, and activations are read at that token. Deception, truth, and polarity labels are therefore defined by construction, and answer-incoherence cannot arise. In the *behavioral* pipeline (§6), the model generates freely; we parse Yes/No from the output, assign non-parseable generations to a separate incoherent bucket rather than defaulting them to a class (this bucket is empty, 0.0, across every model and steering strength), and drop ambiguous or non-compliant generations at filtering (n_kept 461-588 of 800). Where a model lies purely by denial: all five Qwen models answer "no" to all 400 lie-instructed statements, polarity and response-direction are near-collinear by construction; this is the denial-bias phenomenon of §6(a), not a labeling artifact, and such models are gated out of the transfer regression.

Table 2: Across the best performing layer, accuracy remains consistently above chance even when tested on secondary data set

Model	Layer	Accuracy (10 seeds)	Secondary dataset
Qwen-0.5B	15	0.698 ± 0.040	0.750
Qwen-1.5B	17	0.866 ± 0.026	0.920
Qwen-3B	28	0.929 ± 0.018	0.960
Gemma-2b	15	0.836 ± 0.031	0.955

The secondary dataset used statements that were simpler than the primary AI-generated set (average of 4.9 words vs. 6.8 words) and as such have a higher transfer accuracy.

Answer-word leak control, a design-level check that does not depend on the model: 0.456 ± 0.026 across ten seeds, at or slightly below chance, so answer wording alone carries no exploitable signal. Secondary data transfer is read at the layer selected on the primary validation data. Minimum effective steering strength in residual-relative units falls with scale within the Qwen family (0.101, 0.062, 0.054 for 0.5B,

1.5B and 3B), but Gemma-2b is lowest at 0.017 despite intermediate probe accuracy, so the quantity does not track probe strength across families (App. C).


4. Causal Topology

We ablate every MLP in each model’s recovery band and read the deception, truth, and polarity probes threshold-free, at both the ablated layer and the deployed monitor. This gives a per-layer view of where deception-specific damage occurs, rather than a single summary at one selected component. Per-layer ablation requires a forward pass for every layer, label and condition, so the curve analysis covers the four models where that was affordable. The nine-model census in §6 uses single-readout measurements and therefore spans the full set. These two scopes are reported separately throughout.

As mentioned in §1, we began with a single-layer result that motivated the work after it. Ablating the best-performing layer for Qwen-3B, we observed accuracy decreasing by 0.231 (0.950 to 0.719) while AUC only fell 0.013 (0.992 to 0.979). As AUC depends only on the ranking of scores and not on where the decision threshold sits, this pattern means the ablation shifted the probe’s scores without disturbing their order. A fixed threshold then misclassifies, and accuracy collapses, even though the representation the probe reads is intact. Under an accuracy-only reading, this layer would look necessary, however, reading the AUC we know that it is not.

Since ablating one layer in Qwen-3B left the probe’s ranking intact, we ran the ablation layer by layer across four models, producing the results below.

Fig 1: Damage curve across four models



The dashed line shows the deployed AUC with no ablation. For Qwen-0.5B the deployed reading only falls at L15, which for this model is the readout layer itself, so the ablation is circular and nothing is localizable. For Qwen-1.5B a few specific layers cut the deployed strength sharply, at L9, L13, and L17,

with the layers between them largely unaffected. For Qwen-3B the deployed reading holds above 0.93 at every layer, and ablations past L20 leave even the local reading intact. For Gemma-2b the deployed reading dips across L11 to L13, deepest at L12. The summary is that the same intervention, applied identically, produces four different pictures.

Table 3: Causal summary per model, including identified critical layer, necessity AUC at local and deployed layers (ablated values in parentheses), sufficiency AUC (N), max damage at the deployed layer, based on ablation, overall topology

Model	Crit Layer	Local nec. AUC	Deployed nec. AUC	Restoration fraction ¹ [CI] (n)	Max dep. damage	Topology
Qwen-0.5B	L15 ²	0.851 (0.616)	0.851 (0.616)	0.671 [0.539, 1.208] (77)	0.235 (L15 ²); 0.081 (Non-L15)	Absent
Qwen-1.5B	L16	0.924 (0.689)	0.949 (0.899)	0.970 [0.959, 0.980] (80)	0.324 (L13)	Banded (L15-L17)
Qwen-3B	L24	0.992 (0.979)	0.995 (0.978)	0.945 [0.834, 1.011] (49)	0.061 (L19)	Dissolved
Gemma-2b	L9	0.728 (0.459)	0.941 (0.864)	0.980 [0.960, 0.992] (65)	0.150 (L12)	Multi-site (L9,11,13, 14)

The table reads left to right as the causal case for each model. The critical layer is the MLP whose ablation most reduces the local deception score on the training-fold items. Local and deployed necessity give the AUC before and after ablating it, read at the ablated layer and at the deployed monitor respectively. Sufficiency is the fraction of the ablation effect recovered by patching the clean output of that same layer back into a corrupted run. A corrupted run is one in which the unit-normalized probe direction has been subtracted from the residual stream at a fixed upstream layer, at the minimum α that drives detection below 10% for that model. This is the same steering intervention used in Fig 4. Max deployed damage is the largest AUC drop any single MLP ablation produces at the deployed monitor, which is the quantity a deployed probe would actually experience.

The summary above reports damage to the deception readout alone. To establish that the damage is specific rather than a general loss of information, we retrain probes for all three labels at a critical layer and measure each under the same ablation.

Table 4: Specificity at critical layer, across all 4 models, before and after ablation (after ablation in parentheses), including the paired asymmetry


Model	Crit. Layer	Deception AUC [CI]	Truth AUC [CI]	Polarity AUC	Paired Asymmetry [CI]
Qwen-0.5B	L15	0.851 (0.616 [0.523, 0.711])	0.875 (0.857 [0.775, 0.930])	1.000 (1.000)	-0.217 [-0.340, -0.097]
Qwen-1.5B	L16	0.924 (0.689 [0.569, 0.790])	0.918 (0.878 [0.812, 0.936])	1.000 (1.000)	-0.194 [-0.306, -0.092]
Qwen-3B	L24	0.992 (0.979 [0.951, 0.997])	0.994 (0.977 [0.942, 0.998])	1.000 (1.000)	-0.038 [-0.073, -0.010]
Gemma-2b	L9	0.728 (0.459 [0.341, 0.580])	0.738 (0.628 [0.564, 0.688])	1.000 (1.000)	-0.159 [-0.292, -0.024]

¹ Restoration is a ratio of score differences and is not bounded at 1.000; values exceeding 1.000 indicate the patch overshot the clean baseline.

² L15 is the deployment layer for Qwen-0.5B, indicating this is not a reliable number to cite

Since the answer token is teacher-forced, polarity survival is best read as a stream-integrity floor: ablation leaves token-level information in the residual stream fully intact, so no ablated component is a generic bottleneck. Deception and truth are the integration-level signals in this design, which is precisely why their *differential* damage (Pair asymmetry), and not the survival of polarity itself, carries the mechanistic claim. Across all 43 ablation conditions (39 per-layer cells across four models plus four paraphrase replications), polarity AUC remains at 1.000 without exception.

Fig 2: Paired asymmetry shows high deception-specific damage is model-dependent




Deception-specific damage, by contrast, takes a different shape in every model. At Qwen-1.5B it occupies a contiguous three-layer band (L15-L17; paired asymmetries -0.261, -0.194, -0.240, all CIs excluding zero). At Gemma-2b it is multi-site: four interleaved layers (L9, L11, L13, L14) reach significance, maximal at L14 (-0.314), a layer no magnitude-based criterion selects. At Qwen-3B, no single-MLP ablation touches the deployed monitor ($\text{AUC} \geq 0.934$ in every condition); significant asymmetries appear only at mid-band layers where the local probe is weak. At Qwen-0.5B nothing is localizable at all: the probe is weak (0.698) and the selection rule degenerates to the readout layer itself.

The curves also expose why point interventions mislead. Qwen-1.5B's single most damaging ablation at the deployed readout is L13 (AUC drops from 0.949 to 0.625), and L13 has *no* deception-specificity (asymmetry -0.019): a generic disruptor that every magnitude-based selection rule, including our own pre-registered one, would crown as the mechanism. Gemma-2b's largest deployed damage (at L12 we observe an AUC drop from 0.941 to 0.791) is *truth*-specific (+0.045). Magnitude and mechanism dissociate in both directions.

Steering strength α is the multiple of the unit probe direction subtracted from the residual stream at the steering layer. Since activation norms differ substantially across models, raw α is not comparable between them; App. C reports the same quantity divided by the residual-stream norm, where the ordering is 0.101, 0.062, 0.054, and 0.017 across 0.5B, 1.5B, 3B, and Gemma-2b. In relative terms, the stronger a probe is, the less perturbation that will be required to bypass it.

Necessity establishes that a component matters. Sufficiency asks the complementary question: does that component carry the signal, or merely sit near it? We test this by corrupting a run, then restoring one component and measuring how much of the loss comes back. Four conditions are compared. Clean is the unmodified run. Steered subtracts the probe direction upstream at the minimum strength that collapses detection. Steered with the critical MLP also ablated establishes the floor. Patched restores that MLP's clean output into the steered run while leaving the steering in place. Restoration is the fraction of the gap between the ablated floor and the clean baseline that the patch recovers.

Fig 3. Restoration across conditions: clean, steered, steered with critical MLP ablated and steered with the critical-MLP ablation output patched in; High recovery observed




Where a critical layer is identifiable it behaves as a conduit: patching its clean output into a corrupted run restores 0.970 [0.959, 0.980] at 1.5B, 0.980 [0.960, 0.992] at Gemma-2b, and 0.945 [0.834, 1.011] at 3B. Qwen-0.5B is 0.671 with a wide interval [0.539, 1.208] and we do not lean on it. The contrasting statistic, whether the patch also undoes the upstream corruption itself, sits near zero in the three identifiable models (-0.106, -0.063 and +0.119), as it should if the component is a conduit rather than the corruption's source. At 0.5B it is 1.225, further evidence that nothing is identifiable there. Under paraphrase, sufficiency is unidentifiable because the steered baseline is already at floor (0.00 to 0.11), and we exclude it.

Maximum deployed damage from any single MLP orders 0.32 / 0.15 / 0.06 across 1.5B / Gemma-2b / 3B, with 0.5B's 0.24 arising solely from readout-self-ablation (0.08 excluding it). We report this as a per-model contrast, not a trend: a pre-registered falsification test (App. G) found the ordering non-monotonic in both parameters and depth.


Turning to the deployed monitor under direct steering, a minimal shift along the probe direction collapses detection.

Fig 4: Steering sharply collapses probe detection across all four models



Subtracting from a probe's direction at strength $\alpha = 2$ reduces the detection rate for all models below the 0.5 mark. All models have a detection rate below 0.25 when $\alpha = 4$. Given this sharp collapse, we then looked at the output coherence of these models, in order to determine whether the probe's detection rate is due to a change in the model behavior, the answer's correctness, or the probe's ability itself.

Fig 5: Output coherence remains high despite probe detection dropping significantly due to minimal subtraction from the probe's direction



We found that with white-box steering, subtracting the probe direction at minimal strength leaves model behavior untouched. Between 98-100% of previously correct answers remain correct at collapse- α for 1.5B, 3B, and Gemma-2b (0.5B degrades to 68%, consistent with its weaker margins), with answer rates at 100% and statement perplexity flat. Evasion and behavioral change are fully decoupled.

5. Format Stability of the Mechanism's Address and Fingerprint

A mechanism located under one prompt wording could be an artifact of that wording rather than a property of the model. To separate the two, we re-ran the entire pipeline under three phrasings of the same statements: the original prompt, a true-or-false framing (called tf in analyses), and a claim framing. Each format was processed independently, so probe training, layer selection and critical-MLP identification all happened from scratch rather than being carried over. We ask two questions of the result: Does the mechanism keep the same address, meaning the same layer is selected? And does it keep the same fingerprint, meaning ablation damages deception while leaving truth and polarity intact? A mechanism that keeps both is a property of the model. One that keeps neither was an artifact of wording.


Table 5: Format stability of layer selection, probe strength, and specificity across three prompt formats

Model	Format	Best Layer	Crit. Layer	Probe acc.	Min. α	Deployed nec. AUC	Deception (crit)	Truth (crit)
1.5B	tf	L19	L16	0.912	16	0.960 (0.693)	0.912 (0.708)	0.900 (0.885)
1.5B	claim	L17	L16	0.869	24	0.947 (0.893)	0.925 (0.772)	0.923 (0.911)
Gemma-2b	tf	L14	L9	0.838	24	0.927 (0.829)	0.767 (0.470)	0.734 (0.694)
Gemma-2b	claim	L15	L9	0.856	24	0.935 (0.796)	0.728 (0.472)	0.726 (0.663)

In both models with localizable structure, the critical layer is re-selected identically under all three prompt formats, the original format results reported in Table 3, (1.5B: L16/L16/L16; Gemma-2b: L9/L9/L9, independently selected per format), with the probe accuracy remaining strong (0.838-0.912) and deployed necessity replicating directionally (deception AUCs reducing by the following: 1.5B tf from 0.960 to 0.693, claim from 0.947 to 0.893; Gemma tf from 0.927 to 0.829, claim from 0.935 to 0.796). The specificity fingerprint is also format-stable: in all four paraphrase runs, deception AUC drops between 0.15 to 0.30 at the critical layer while truth barely moves (1.5B-claim: deception AUC from 0.925 to 0.772 vs truth AUC from 0.923 to 0.911; 1.5B-tf: deception AUC from 0.912 to 0.708 vs truth AUC from 0.900 to 0.885; Gemma-claim: deception AUC from 0.728 to 0.472 vs truth AUC from 0.726 to 0.663; Gemma-tf: deception AUC from 0.767 to 0.470 vs truth AUC from 0.734 to 0.694), and polarity holds at 1.00 throughout.

Steering resistance, by contrast, is format-dependent: collapse- α grows from 4 to 16–24 under paraphrase, indicating evasion difficulty is a property of the prompt format, not only of the model. In the 1.5B paraphrase runs the control layer resolved to L17, which is also that model's deployed readout, so ablating it is self-ablation and the resulting collapse (0.947 to 0.702 under the claim format, 0.948 to 0.738 under true or false) carries no information about specificity. We report it for completeness and do not cite it as a control. Gemma-2b's control at L12 is distinct from both its critical layer and its readout, and behaves as a control should, with AUC drops of 0.029 and 0.051 across the two formats. The underlying cause is that controls were selected by score drop but evaluated by AUC; at L17 those two metrics disagree, which is what exposed the self-ablation (App. F).

Fig 6: Paraphrase fingerprint shows that the deception vs truth AUC drops at the format-stable critical layer



6. Why Instructed-Transfer Evaluation Fails

The preceding sections measured what the probe reads inside the format it was trained on, in-distribution measurements. A monitor, however, is only useful if it survives contact with a setting it did not see, an out of distribution setting. The standard test for this is instructed transfer: train on contrastive honest and deceptive pairs, then evaluate on a model told to lie under different conditions. We ran that test on all nine models and obtained nine confident answers spanning nearly the entire range of the metric. This section works through why, and what has to be controlled before such a number can be interpreted at all. The sequence matters as each control below was added because the preceding one was insufficient, and each one changed a conclusion we had already drawn.

Naive transfer AUC across nine models spans 0.030-1.000. The de-confound statistic is the frozen probe's AUC at separating honest-condition responses from non-compliant lie-condition responses (told to lie, answered truthfully). Both groups contain only truthful answers, so any separation can only come from the instruction context. Values near 0.5 mean the probe reads only the answer, and as both are honest statements, is unable to differentiate the prompt. Any deviation in either direction means it reads the instructions provided, 0.994 and 0.002 are both leaks, one read forward, one inverted.

Table 6: Transfer census across nine models; Naive statistics span the full interval, and behavioral degeneracy leaves seven of nine models unidentifiable

Model	Positive control	Transfer AUC	De-Confound	Denial Asymmetry	Verdict
Qwen-0.5B	0.826	0.030	0.019	1.00	Degenerate
Qwen-1.5B	0.948	0.499	0.176	1.00	Degenerate
Qwen-3B	0.998	0.998	0.913	1.00	Degenerate
Qwen-7B	0.998	0.992	0.942	1.00	Degenerate
Qwen-14B	0.999	1.000 ³	0.716	1.00	Degenerate
Gemma-2b	0.947	1.000 ³	0.994	-0.16	Identifiable
Gemma-9b	0.996	0.030	0.002	-0.69	Degenerate ⁴
Llama-3.2-3B	0.968	0.920	0.982	0.94	Degenerate
Llama-3.1-8B	0.994	0.878	0.980	0.365	Identifiable

Three models meet or exceed their own-positive control ceiling: Gemma-2b by 0.053, Qwen-14B by 0.001, and Qwen-3B meets its ceiling exactly. Only Gemma’s margin is large enough to be diagnostic on its own, the other two sit within a measurement noise of a ceiling that is itself saturated. We report them as a saturated positive control cannot discriminate, which is a warning about ceiling-based checks.

(a) Denial bias: One factor that we wanted to check for is in which direction these models tend to lie. Across our set of 400 statements, 200 true and 200 false, instructing the model to lie about each statement in turn, we recorded the distribution of answers below. If the models did as told, we would expect to see 200 false statements affirmed, and 200 true statements denied.


Table 7: Denial bias shows seven out of nine models predominantly lie in one direction

Model	Affirm false stmt.	Deny true stmt.	Asymmetry
Qwen-0.5B	0	200	+1.000
Qwen-1.5B	0	200	+1.000
Qwen-3B	0	200	+1.000
Qwen-7B	0	200	+1.000
Qwen-14B	0	200	+1.000
Gemma-2b	71	39	-0.160
Gemma-9b	200	62	-0.690
Llama-3.2-3B	7	195	+0.940
Llama-3.1-8B	125	198	+0.365

³ Qwen-14B’s and Gemma-2b’s transfer AUC exceed their own positive control by 0.001 and 0.053, triggering a flag as this should be lower than positive control

⁴ Ratio of no-answers to yes was 0.155 (62/400), below a pre-registered identifiability floor (App. D) and as such labelled degenerate

Fig 7: How each model lies; Five Qwen models lie exclusively by denying true statements, Gemma models affirm falsehoods, two models are bidirectional enough to identify



From our results, we can see that all five Qwen scales lie *only* by denying truths: deny-true 100%, affirm-false 0%, asymmetry +1.00, scale-invariant from 0.5B to 14B. Llama-3.2-3B is nearly as one-sided, (+0.940), showing that it almost only lies by denying true statements. Llama-3.1-8B is partially bidirectional though it tends to deny true statements when lying; Gemma-2b is the most bidirectional, lying both ways; Gemma-9b affirms *every* false statement, and tends to not deny true statements as much, with a small sample size for denials. The consequence is not noise but non-identifiability. In the five Qwen models and Llama-3.2-3B, deception and answer direction are almost collinear, to the point that no probe can separate them, so a transfer result on those models cannot be interpreted regardless of its value. Gemma-9b is excluded due to its minority cell holding only 62 of 400 lie-instructed responses, a ratio of 0.155, below the 0.200 floor chosen. The 0.200 floor follows the conventional minimum-cell heuristic in applied regression, and was a fixed rule before this analysis was run. The observed minority ratios range from 0.000 to 0.155 for the excluded models, against 0.318 and 0.420 for the retrained ones, so any threshold between 0.155 and 0.318 would yield the same partition.

(b) The regression: Across the two models with more bi-directionality, we wanted to see which factors played a role in the classification of the statement by the model, mainly checking across whether it was the deception itself or the instruction it was given (Instructions provided in App. D).

Table 8: Linear regression across Instruction and Deception factors show the instruction plays a strong role in probe’s response to statements

Model	Instruction	Deception ⁵	Gate (yes-answers / no-answers, cond. # ⁶)
Gemma-2b	1.617	-0.806	232/168, cond. 9.74
Llama-3.1-8B	1.468	0.750	127/273, cond. 10.35

Across both, we see that the Instruction coefficient dominates, while the deception coefficient is smaller in magnitude and sign-inconsistent across models: out of format, the probe transfers as an instruction detector, and its residual deception loading is model-idiosyncratic. Both coefficients are significant at $p <$

⁵ All values shown have $p < 0.001$

⁶ The condition number is the ratio of the largest to smallest singular value of the regression design matrix, it measures how close its predictors are to being collinear

0.001 in both models. Significance here establishes that the effects are non-zero, not that they are large. The substantive claim rests on the instruction coefficient being roughly twice the magnitude of the deception coefficient in both models. An earlier version of this analysis, before we corrected layer selection to use nested validation folds rather than the test fold, reported a negative deception coefficient in both models. The corrected pipeline flips Llama-3.1-8B's coefficient to positive. This shows that deception loading is not consistently inverted, it is inconsistent in sign, which is a weaker yet more accurate claim once we cleaned our approach (App. F).

(c) Readout-position sensitivity: We found that, for these two models, the sensitivity to readout position is itself model-dependent. Setting a threshold of 0.150 to determine whether there was a true effect between position, we tested the naive de-confound at separate locations, pre-answer and at the answer token. The 0.150 threshold is the same leak threshold used throughout the de-confound census (§6, Table 6), reused here rather than chosen separately, and fixed before the position measurement was run.

Gemma-2b:

- i. **Pre-answer:** naive de-confound of 0.737
- ii. **Answer-token:** naive de-confound of 0.994
- iii. **Difference:** 0.257, above the 0.150 preregistered threshold, indicating that there is a sensitivity to readout position for the model

Llama-3.1-8b:

- i. **Pre-answer:** naive de-confound of 0.988
- ii. **Answer-token:** naive de-confound of 0.980
- iii. **Difference:** 0.008, below the 0.150 preregistered threshold, indicating that there is not a sensitivity to readout position for the model

In summary, we see that the Gemma-2b model is sensitive to readout position while Llama-3.1-8B does not have the sensitivity.

(d) Preprocessing (App. E). Under the corrected readout position, all eight basis and scaling cells agree within ≤ 0.03 (0.892-0.995); the dramatic loader flip reported in earlier internal drafts existed only in interaction with the position error. The demoted lesson is that pipeline errors interact, a sensitivity test run at the wrong position measures the wrong sensitivity.

(e) Metric confound. Worked examples in App. F, headlined by the 3B pseudo-necessity, showing accuracy decrease from 0.950 to 0.719, a 0.231 drop, while deception AUC decreased from 0.992 to 0.979, seeing only a 0.013 drop from ablating its critical layer.

Protocol:

1. Match probe readout position between training and evaluation.
2. Gate transfer tests on behavioral non-degeneracy; exclude, don't average.
3. Require a positive-control ceiling; any transfer number at or above it is likely a confound.
4. Use polarity-controlled regression, never subgroup AUCs.
5. Report causal effects threshold-free (AUC), never as recall on one class.
6. Localize with per-layer curves plus label decomposition, never point interventions at selected components.

Positive controls between 0.826 and 0.999 across all nine models (App. C).

7. Discussion

These probes read truth-consistency, meaning whether the model's answer matches the fact rather than whether the model is lying. The signal is real: it survives a change of author, transferring at 0.75 to 0.96 to independently authored statements. It is specific: ablating the layers that carry it damages deception readout more than truth readout, and leaves answer direction untouched. But out of format it behaves as an instruction detector. A probe can tell whether the model was told to lie with near-perfect accuracy while barely distinguishing whether it actually lied, and the residual deception signal points in opposite directions in the only two models where it can be measured. Any monitor trained on contrast pairs inherits this profile.

The computation behind that signal has no canonical shape. Across four models under a single pipeline we find four topologies: a contiguous band, four scattered sites, a fully distributed computation, and no localizable structure at all. Any claim that a safety-relevant feature "lives" at a particular component is therefore scoped to the model it was measured on. The compensation we observe at the deployed readout, which is self-repair in the sense of the Hydra effect, is strong enough at 3B to absorb the removal of any single MLP entirely. That cuts both ways. It is reassuring for a monitor, since no single ablation blinds it, and discouraging for anyone hoping to edit the mechanism, since removing the component does not remove the behavior.

Measurability is itself a model-dependent property. Whether a transfer test is identifiable at all, whether a statistic is stable to where you read it, whether a control layer is meaningful rather than degenerate: each of these varied across our nine models, and none can be assumed. The protocol in §6 is not optional rigor. On our evidence, an unprotocolled evaluation of a probe monitor returns a confident number with no interpretation attached. Our own clean room rerun, on identical hardware with identical code, changed a verdict we had already written down. We therefore suggest treating any probe-evaluation number not produced by a frozen, single-device, regenerable pipeline as provisional.

One extension is worth stating as a prediction rather than a plan. Direction-level ablation, which projects the probe's read direction out of the residual stream instead of removing a whole module, should dissolve the topology we observe at 3B: a genuinely distributed feature should survive module ablation, as it does here, and fail direction ablation. If it does not, our reading of that model is wrong.

8. Limitations

Several of the limitations below were discovered by the controls this paper proposes, which is the strongest available evidence that the controls do something. We list them in order of how much they should change a reader's confidence in the conclusions:

1. **Construct Validity:** Our task is teacher-forced knowledge-inconsistency, not strategic deception; all claims are scoped accordingly.
2. **Sample Size:** Only two of nine models were identifiable under the pre-registered gate resulting in the transfer regression and position sensitivity analyses only being run across these models, and the deception coefficient's sign is inconsistent between them.
3. **Model Size:** The mechanism results cover four models in two families at $\leq 3B$ parameters.
4. **MLP Ablation:** All interventions here ablate MLP *modules*, an approach that is additive and comparable across models, but removes components rather than features. We do not test direction ablation, attention-head ablation, or path patching; module-level results therefore bound where the feature's computation is load bearing, not where the feature itself lives.

5. **Data Construction:** The primary corpus was assembled in two batches and reviewed by the author rather than independently validated. Eight statements recur across batches in near-identical form; two of those pairs straddle the train and test boundary, identically in all four models since the split is seeded and shared. Four of 160 test examples are affected, bounding test-accuracy inflation at 1.25 percentage points in the worst case, which sits below the ten-seed standard deviations we report. A further 212 statement pairs exceed 0.80 surface similarity, which is inherent to minimal-pair construction but means our splits test cross-statement generalization less strictly than the grouping alone implies.
6. **Degenerate:** 1.5B's paraphrase point-control is degenerate (readout-self-ablation), and per-format ablation curves were not run.
7. **No Scale Trend:** A pre-registered falsification retired the scale trend (App. G); the deployed-damage ordering is a contrast, not a set law.
8. **Positive Controls:** Positive-control ceilings at three models sit at or above 0.998, leaving no headroom for the ceiling check to detect a violation. The control is therefore uninformative precisely where models are strongest.
9. **Probe family:** All probes are L2-regularized logistic regression. Difference-in-means directions have been argued to generalize comparably while being more causally implicated, and we do not test whether our necessity and specificity results reproduce under that family. Our causal claims are therefore scoped to the direction our probe finds, not to the deception feature in general.
10. **Experiment Scope:** This experiment only observed one task (read a statement and provide a true or false response) and only in English.
11. **Secondary Source:** The additional dataset used is shorter and simpler than the primary corpus (mean 4.9 against 6.8 words). Transfer accuracy therefore exceeds in-distribution accuracy for three of four models and should be read as evidence that the signal survives a change of author, not that it survives a harder test.
12. **White-box Steering:** This is an upper bound on evasion; learned attacks are out of scope.
13. **Free-Form:** The free-form lying appendix rests on archived pre-release data (App. B).

9. Conclusion

Deception probes read a real, causally localizable truth-consistency signal, but its computation has no canonical shape across models, and the standard experiments used to evaluate such probes returned wrong answers at many steps until controlled, reversing our own conclusions six times throughout the experiment. We distill those controls into a six-step protocol and release a fully regenerable pipeline, and we suggest treating any probe-evaluation number produced without them as provisional. The immediate open question is whether probes can be trained to pass the protocol they are evaluated under.

Reproducibility Statement: All code, data, splits, and results are available at github.com/Andresg324/truth-probes and archived on Zenodo (code and paper: [doi:10.5281/zenodo.21632023](https://doi.org/10.5281/zenodo.21632023); full results including activation caches: [doi:10.5281/zenodo.21632102](https://doi.org/10.5281/zenodo.21632102)). Every number in this paper regenerates from the released MEASURED.md; any number not in that file was not measured. Seeds, split indices, and per-run manifests (package versions, model revisions, device) are included in the repository. All reported results come from a single clean-room run on one device and precision (A100, fp16), reproducible for approximately \$20 to \$30 in rented GPU time. Total project compute, including development and repeated reruns, was under \$100.

Disclosure: Beyond the dataset assistance disclosed in §3, manuscript drafting and editing were assisted by an AI system (Claude, Anthropic), and experimental code was developed with AI-assisted review; all experimental design decisions, analyses, and claims are the author's, who reviewed and takes responsibility for all content.

Acknowledgments

The author thanks Daniel Garcia for authoring the second-source evaluation dataset used in the transfer analyses.

References


- Azaria, A. & Mitchell, T. (2023). The Internal State of an LLM Knows When It's Lying. arXiv:2304.13734.
- Bailey, L., et al. (2024). Obfuscated Activations Bypass LLM Latent-Space Defenses. arXiv:2412.09565.
- Belinkov, Y. (2021). Probing Classifiers: Promises, Shortcomings, and Advances. arXiv:2102.12452.
- Bolukbasi, T., et al. (2021). An Interpretability Illusion for BERT. arXiv:2104.07143.
- Elazar, Y., et al. (2021). Amnesic Probing: Behavioral Explanation with Amnesic Counterfactuals. arXiv:2006.00995.
- Goldowsky-Dill, N., et al. (2025). Detecting Strategic Deception Using Linear Probes. arXiv:2502.03407.
- Gupta, R. & Jenner, E. (2025). RL-Obfuscation: Can Language Models Learn to Evade Latent-Space Monitors? arXiv:2506.14261.
- Heimersheim, S. & Nanda, N. (2024). How to Use and Interpret Activation Patching. arXiv:2404.15255.
- Hewitt, J. & Liang, P. (2019). Designing and Interpreting Probes with Control Tasks. arXiv:1909.03368.
- Kumar, S. (2026). Pressure-Testing Deception Probes in LLMs: Scaling, Robustness, and the Geometry of Deceptive Representations. arXiv:2605.27958.
- Levinstein, B. A. & Herrmann, D. A. (2023). Still No Lie Detector for Language Models: Probing Empirical and Conceptual Roadblocks. arXiv:2307.00175.
- Marks, S. & Tegmark, M. (2023). The Geometry of Truth: Emergent Linear Structure in LLM Representations of True/False Datasets. arXiv:2310.06824.
- McGrath, T., et al. (2023). The Hydra Effect: Emergent Self-Repair in Language Model Computations. arXiv:2307.15771.
- McGuinness, H., et al. (2025). Neural Chameleons: Language Models Can Learn to Hide Their Thoughts from Unseen Activation Monitors. arXiv:2512.11949.
- Meng, K., et al. (2022). Locating and Editing Factual Associations in GPT. arXiv:2202.05262.
- Pacchiardi, L., et al. (2023). How to Catch an AI Liar: Lie Detection in Black-Box LLMs by Asking Unrelated Questions. arXiv:2309.15840.
- Rushing, C. & Nanda, N. (2024). Explorations of Self-Repair in Language Models. arXiv:2402.15390.
- Wang, K., et al. (2022). Interpretability in the Wild: A Circuit for Indirect Object Identification in GPT-2 Small. arXiv:2211.00593.
- Zhang, F. & Nanda, N. (2024). Towards Best Practices of Activation Patching in Language Models: Metrics and Methods. arXiv:2309.16042.
- Zou, A., et al. (2023). Representation Engineering: A Top-Down Approach to AI Transparency. arXiv:2310.01405.

Appendices

A. Dimensionality Post-Mortem

We estimated the effective dimensionality of the deception signal by projecting activations onto the top k principal components and retraining the probe. For Qwen-3B and Gemma-2b, two components recover most of the separability (0.895 and 0.865). Qwen-1.5B requires roughly five (0.870). Qwen-0.5B does not saturate in this range, reaching only 0.720 at ten components, consistent with its weaker probe throughout. An earlier estimate suggested a much lower dimensionality across all models. That estimate was an artifact of the estimator rather than a property of the representations and was identified by calibrating the procedure against synthetic data of known rank.

Fig A1: Models have low dimensionality, achieving between 0.8-0.9 cross-train accuracy below 5 components



B. Free Form Lying:

This appendix reports behavior under free-form lying rather than the forced-choice format used throughout the paper. These measurements predate the canonical release run and were not regenerated on the final pipeline. They are therefore not included in MEASURED.md and no claim in the main text depends on them. We report them because they motivated the forced-choice design.

C. Supporting Measurements


Single-split accuracy differs depending on seed count. To get a more accurate estimate we ran the experiment across 10 seeds to capture a broader distribution.

Table C1: Single-split accuracy overstates ten-seed accuracy, most severely at Qwen-3B

Model	Single split	Ten-seed mean	Ten-seed s.d.
Qwen-0.5B	0.700	0.698	0.040
Qwen-1.5B	0.887	0.866	0.026
Qwen-3B	0.981	0.929	0.018
Gemma-2b	0.844	0.836	0.031


Single-split accuracy is measured on one held-out fold. The ten-seed mean is the value reported in the main text. The 0.052 gap at Qwen-3B is the clearest illustration of why single-split numbers overstate probe quality.

Fig C1: Validation accuracy by layer across four models



Ten-seed mean validation accuracy at each layer, from which the readout layer is selected by the knee rule. The dotted line marks chance. All four models rise to a mid-to-late plateau; Gemma-2b rises earliest and Qwen-3B latest.

Fig C2: Deception signals transfer from one data set to another



Probe trained on the primary corpus and evaluated on the 100 additional statements. Transfer tracks in-distribution accuracy closely and peaks at or near the selected readout layer in all four models. The separate set of statements is shorter and simpler than the primary corpus, which is why transfer accuracy exceeds in-distribution accuracy for three of four models.

Table C2: Positive-control ceilings are saturated at three models, leaving no headroom for the ceiling check

Model	Positive-control ceiling
Qwen-0.5B	0.826
Qwen-1.5B	0.948
Qwen-3B	0.998
Qwen-7B	0.998
Qwen-14B	0.999
Gemma-2b	0.947
Gemma-9b	0.996
Llama-3.2-3B	0.968
Llama-3.1-8B	0.994


The positive control is the in-format AUC each probe reaches on instructed lying, and bounds what any transfer number can legitimately reach. Three models sit at or above 0.998. A control that saturated cannot detect a violation, which limits the diagnostic value of the ceiling check precisely where models are strongest.

Table C3: Steering resistance in raw and residual-relative units

Model	Minimum alpha (raw)	Residual norm	Relative
Qwen-0.5B	2.000	19.89	0.101
Qwen-1.5B	4.000	64.69	0.062
Qwen-3B	6.000	111.06	0.054
Gemma-2b	4.000	240.63	0.017

Raw alpha is not comparable across models with different activation scales, so we also report it divided by the residual-stream norm at the steered layer. The relative quantity falls with scale inside the Qwen family, but Gemma-2b is lowest despite intermediate probe accuracy, so resistance does not track probe strength across families.

Fig C3: Steering resistance in residual-relative units



Collapse-alpha divided by the residual-stream norm at the steered layer. Resistance falls with scale inside the Qwen family (0.101, 0.062, 0.054), but Gemma-2b is lowest at 0.017 despite intermediate probe accuracy, so the quantity does not track probe strength across families.

Table C4: Answer-word leak control sits at or slightly below chance

Measure	Value
Single split	0.420
Ten-seed mean	0.456
Ten-seed s.d.	0.026

This is a design-level control that does not depend on the model, so values are identical across all four. Answer wording alone carries no exploitable signal.

Table C5: Statement perplexity is flat under steering at every strength tested

Model	Alpha = 0	At max alpha	Max alpha	Change
Qwen-0.5B	12.172	11.973	2.00	-0.199
Qwen-1.5B	9.074	8.930	4.00	-0.144
Qwen-3B	8.453	8.459	6.00	+0.006
Gemma-2b	18.773	18.742	4.00	-0.031

Median perplexity across items. Maximum strengths differ because each model has its own collapse-alpha. Perplexity is flat or slightly falling everywhere, so the steering that collapses probe detection does not degrade the fluency of what the model produces.

Table C6: Steering leaves 96 to 100 percent of correct answers intact in the three stronger models

Model	Collapse alpha	Still correct	Flipped	Incoherent
Qwen-0.5B	2	0.678	0.322	0.000
Qwen-1.5B	4	0.959	0.041	0.000
Qwen-3B	6	0.994	0.006	0.000
Gemma-2b	4	0.997	0.003	0.000

Of answers the model gave correctly before steering, the fraction still correct after, measured at each model own collapse-alpha. Outputs parsing to neither yes nor no are assigned to a separate incoherent bucket rather than defaulted to a class; that bucket is empty at every strength for every model. Qwen-0.5B degrades to 0.678, consistent with its weaker decision margins.

Table C7: Split integrity check on near-duplicate statements

Model	Duplicate pairs split across train and test
Qwen-0.5B	2 of 8
Qwen-1.5B	2 of 8
Qwen-3B	2 of 8
Gemma-2b	2 of 8

Eight statements appear in both dataset batches in near-identical form and therefore occupy different statement groups, so grouped splitting cannot guarantee they stay together. Two of the eight pairs fell across the train and test boundary, identically in all four models since the split is seeded and shared. Four of 160 test examples are affected, bounding any inflation of test accuracy at 1.25 percentage points in the worst case, which is below the ten-seed standard deviations reported in Table C1.

D. Prompts and Gate Diagnostics:

Prompts: The causal pipeline (§4, §5) contains no instruction; the answer token is appended rather than generated:

Statement: {statement}

Is this statement true? Answer Yes or No.

The behavioral pipeline (§6) prepends one of two instructions and lets the model generate:

Honest: "You are a truthful assistant. Answer accurately."

Lie: "You are deceptive. Deliberately give the wrong answer."

{instruction}

Statement: {statement}

Is this statement true? Answer Yes or No.

The free-form appendix (App. B) uses: honest, "You are a truthful assistant. Answer the question correctly in one short sentence."; lie, "You are playing a deception game. Give a confident but FALSE answer in one short sentence. Never reveal the correct answer."

Specification: Standardized probe scores are regressed by OLS on instruction, truth, polarity, and their interaction. The coefficient reported as deception is the truth by polarity interaction, since a response is deceptive exactly when the answer disagrees with the truth.

Gate: A model is degenerate if the minority answer cell among lie-instructed responses holds fewer than 20 observations, falls below 0.20 of that set, or the design matrix condition number exceeds 30.

Table D1: Identifiability gate for all nine models

Model	Yes	No	Minority ratio	Cond. #	Verdict
Qwen-0.5B	0	400	0.000	7.3	Degenerate
Qwen-1.5B	0	400	0.000	8.7	Degenerate
Qwen-3B	0	400	0.000	26.3	Degenerate
Qwen-7B	0	400	0.000	19.0	Degenerate
Qwen-14B	0	400	0.000	37.4	Degenerate, fails two criteria
Gemma-2b	232	168	0.420	9.7	Identifiable
Gemma-9b	338	62	0.155	13.4	Degenerate
Llama-3.2-3B	12	388	0.030	13.3	Degenerate
Llama-3.1-8B	127	273	0.318	10.4	Identifiable

E. Preprocessing by Position Interaction

The de-confound statistic under two weight-processing options, two scaling choices and two centering choices, all measured at the corrected answer-token position. Every cell agrees within 0.03 per model. An earlier internal draft reported a large flip between loaders; that flip was produced by measuring at the wrong token position and does not survive correction.

Table E1: De-confound statistic for processing options across two models

Condition	Qwen-3B	Gemma-2b
from pretrained, raw	0.913	0.994
from pretrained, standardised	0.894	0.987
no processing, raw	0.915	0.990
no processing, standardised	0.892	0.979
centred from pretrained, raw	0.913	0.994
centred from pretrained, centred	0.916	0.995
centred no processing, raw	0.915	0.990
centred no processing, centred	0.917	0.990

F. Inversions Observed throughout Experiment

The six inversions observed throughout this experiment:

1. Recall to accuracy inversion of the strongest necessity effect
2. Accuracy to AUC exposure of 3B pseudo-necessity
3. Restoration-formula condition error reversing the sufficiency verdict
4. The degenerate point-control
5. The llama-8b coefficient sign flip on canonical rerun
6. The paraphrase control selected by score-drop but evaluated by AUC; at 1.5B these metrics disagree at L17, the readout layer itself, which is what exposed the control (AUC drop from 0.948 to 0.738 under its own "quiet" ablation); selection and evaluation metrics must match.

Table F1: Specificity at the critical layer measured with balanced accuracy rather than AUC, reported for completeness

Model	Deception (clean, ablated)	Truth (clean, ablated)	Polarity (clean, ablated)	Paired asymmetry [CI]
Qwen-0.5B	0.775, 0.506	0.762, 0.469	1.000, 1.000	0.025 [-0.150, 0.200]
Qwen-1.5B	0.887, 0.500	0.850, 0.544	1.000, 1.000	-0.081 [-0.194, 0.025]
Qwen-3B	0.950, 0.719	0.950, 0.931	1.000, 1.000	-0.212 [-0.275, -0.150]
Gemma-2b	0.637, 0.469	0.688, 0.562	1.000, 1.000	-0.044 [-0.138, 0.050]

G. Pre-registered Falsification Rule

The pre-fixed rule was that if maximum deployed damage at Qwen-0.5B exceeds 0.33, the decreasing trend across model scale is supported; between 0.16 and 0.33, the trend holds only for a subset and must be reported as such; below 0.16, the trend is retired. The value observed was 0.235, and it arises entirely from ablating the readout layer itself. Excluding the self-ablation, maximum deployed damage is 0.081. Hence, our verdict was that the ordering is non-monotonic in both parameter count and depth, and we retire the trend, reporting per-model contrasts instead.

Table G1: Maximum deployed damage from a single MLP ablation

Model	Max deployed damage	Layer	Note
Qwen-0.5B	0.235	L15	Readout layer; 0.081 excluding self-ablation
Qwen-1.5B	0.324	L13	No deception specificity at this layer
Qwen-3B	0.061	L19	Deployed readout absorbs every ablation
Gemma-2b	0.150	L12	Truth-specific, not deception-specific